

ML Project - Comparison of Machine Learning Models for Workout Analysis

```
pip install pandas numpy matplotlib seaborn scikit-learn
```

```
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages (2.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (2.0.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (3.10.0)
Requirement already satisfied: seaborn in /usr/local/lib/python3.11/dist-packages (0.13.2)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.11/dist-packages (1.6.1)
Requirement already satisfied: python-dateutil<=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas) (2025.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (4.56.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (1.4.8)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (24.2)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (11.1.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib) (3.2.3)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.14.1)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
```

Loading the Dataset onto Colab

```
import pandas as pd

file_path = "RecGym.csv"
data = pd.read_csv(file_path, encoding='utf-8')

print("First 10 rows of the dataset:")
print(data.head(10).to_string())

print("\nDataset Info:")
print(data.info())

print("\nSummary Statistics:")
print(data.describe())

print("\nMissing Values:")
print(data.isnull().sum())
```

```
8      1  wrist      1.0  0.500625  0.499625  0.501250  0.499025  0.500275  0.500119  0.501085  Null
9      1  wrist      1.0  0.500000  0.498625  0.502250  0.502381  0.500231  0.499962  0.499743  Null
```

```
Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31155 entries, 0 to 31154
Data columns (total 11 columns):
 #   Column      Non-Null Count  Dtype
---  ---
 0   Subject    31155 non-null  int64
 1   Position    31155 non-null  object
 2   Session     31154 non-null  float64
 3   A_x         31154 non-null  float64
 4   A_y         31154 non-null  float64
 5   A_z         31154 non-null  float64
 6   G_x         31154 non-null  float64
 7   G_y         31154 non-null  float64
 8   G_z         31154 non-null  float64
 9   C_1         31154 non-null  float64
10  Workout     31154 non-null  object
dtypes: float64(8), int64(1), object(2)
memory usage: 2.6+ MB
None
```

```
Summary Statistics:
      Subject  Session      A_x      A_y      A_z  \
count  31155.0  31154.0  31154.000000  31154.000000  31154.000000
mean      1.0      1.0      0.502240      0.498696      0.498309
std      0.0      0.0      0.015989      0.019553      0.017103
```

min	0.102094	0.290400	0.119419	0.000000
25%	0.496406	0.497894	0.496331	0.498615
50%	0.499962	0.500081	0.499919	0.500189
75%	0.503594	0.503162	0.502878	0.501841
max	0.858631	0.669844	0.726325	1.000000

```
Missing Values:
Subject      0
Position     0
Session      1
A_x          1
A_y          1
A_z          1
G_x          1
G_y          1
G_z          1
C_1          1
Workout      1
dtype: int64
```

Checking the Dataset for missing values

```
import pandas as pd

file_path = "RecGym.csv"
data = pd.read_csv(file_path, encoding='utf-8')

missing_values = data.isnull().sum()

print("Missing Values in Each Column:")
print(missing_values[missing_values > 0])
if data.isnull().values.any():
    print("\nThe dataset contains missing values.")
else:
    print("\nNo missing values found in the dataset.")
```

```
Missing Values in Each Column:
C_1      1
Workout  1
dtype: int64

The dataset contains missing values.
```

✓ Data Preprocessing Steps

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split

file_path = "RecGym.csv"
data = pd.read_csv(file_path, encoding='utf-8')
data.dropna(inplace=True)

sensor_columns = ["A_x", "A_y", "A_z", "G_x", "G_y", "G_z", "C_1"]

data[sensor_columns] = data[sensor_columns].apply(pd.to_numeric, errors='coerce')
data.dropna(inplace=True)
scaler = StandardScaler()
data[sensor_columns] = scaler.fit_transform(data[sensor_columns])

pca = PCA(n_components=5)
data_pca = pca.fit_transform(data[sensor_columns])

label_encoder = LabelEncoder()
data["Workout"] = label_encoder.fit_transform(data["Workout"])

X = pd.DataFrame(data_pca)
y = data["Workout"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

print("Preprocessing completed with feature reduction, augmentation, and real-time adaptation!")
print("\nFirst 10 Rows of Preprocessed Data (PCA-Reduced Features):")
print(X_train.head(10))

print("\nEncoded Workout Labels (First 10 Rows):")
print(y_train.head(10))
```

```
print("\nShape of Training and Testing Sets:")
print(f"X_train: {X_train.shape}, X_test: {X_test.shape}")
print(f"y_train: {y_train.shape}, y_test: {y_test.shape}")
```

➦ Preprocessing completed with feature reduction, augmentation, and real-time adaptation!

First 10 Rows of Preprocessed Data (PCA-Reduced Features):

	0	1	2	3	4
721324	0.026924	0.750032	-0.217873	0.280247	-0.055111
264856	0.074519	0.506825	0.332641	-0.743139	2.244520
479048	0.037262	0.176077	-0.036703	0.004328	-0.042720
370383	0.149144	0.059629	-0.100245	-0.079551	0.128406
542760	0.521736	0.453749	0.081733	-0.129390	-0.151991
667097	0.715391	-0.052713	-1.664123	1.298820	-0.664162
27970	0.033977	0.184751	0.005338	0.037680	-0.040637
415685	0.120025	0.165496	-0.475122	0.353599	0.181185
439624	0.046048	-0.425226	1.070706	0.560121	-0.449708
572047	0.184411	0.426513	-0.068706	0.163932	-0.627696

Encoded Workout Labels (First 10 Rows):

721324	2
264856	5
479048	5
370383	5
542760	6
667097	11
27970	5
415685	5
439624	5
572047	5

Name: Workout, dtype: int64

```
Shape of Training and Testing Sets:
X_train: (655356, 5), X_test: (163839, 5)
y_train: (655356,), y_test: (163839,)
```

Reshaping data for ResCNN

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Input, Conv1D, BatchNormalization, ReLU, Add, GlobalAveragePooling1D, Dense
from tensorflow.keras.utils import to_categorical
from sklearn.model_selection import train_test_split
```

```
num_classes = len(np.unique(y_train))
y_train_cat = to_categorical(y_train, num_classes)
y_test_cat = to_categorical(y_test, num_classes)
```

```
time_steps = 20
num_features = X_train.shape[1]
```

```
def create_time_series_data(X, y, time_steps):
    sequences, labels = [], []
    for i in range(len(X) - time_steps):
        sequences.append(X.iloc[i : i + time_steps].values)
        labels.append(y[i + time_steps])
    return np.array(sequences), np.array(labels)
```

```
X_train_seq, y_train_seq = create_time_series_data(X_train, y_train_cat, time_steps)
X_test_seq, y_test_seq = create_time_series_data(X_test, y_test_cat, time_steps)
```

```
print(f"Reshaped data for CNN: {X_train_seq.shape}")
```

➦ Reshaped data for CNN: (655336, 20, 5)

Building ResCNN Model

```
def build_rescnn(input_shape, num_classes):
    inputs = Input(shape=input_shape)

    x = Conv1D(filters=64, kernel_size=3, padding="same")(inputs)
    x = BatchNormalization()(x)
    x = ReLU()(x)

    def residual_block(x, filters=64):
        res = Conv1D(filters, kernel_size=3, padding="same")(x)
```

```

    res = BatchNormalization()(res)
    res = ReLU()(res)
    res = Conv1D(filters, kernel_size=3, padding="same")(res)
    res = BatchNormalization()(res)
    x = Add()([x, res]) # Residual connection
    x = ReLU()(x)
    return x

x = residual_block(x)
x = residual_block(x)

x = GlobalAveragePooling1D()(x)
outputs = Dense(num_classes, activation="softmax")(x)

model = Model(inputs, outputs)
return model

rescnn_model = build_rescnn(input_shape=(time_steps, num_features), num_classes=num_classes)
rescnn_model.compile(optimizer="adam", loss="categorical_crossentropy", metrics=["accuracy"])

rescnn_model.summary()

```

Model: "functional"

Layer (type)	Output Shape	Param #	Connected to
input_layer (InputLayer)	(None, 20, 5)	0	-
conv1d (Conv1D)	(None, 20, 64)	1,024	input_layer[0][0]
batch_normalization (BatchNormalization)	(None, 20, 64)	256	conv1d[0][0]
re_lu (ReLU)	(None, 20, 64)	0	batch_normalization[0...
conv1d_1 (Conv1D)	(None, 20, 64)	12,352	re_lu[0][0]
batch_normalization_1 (BatchNormalization)	(None, 20, 64)	256	conv1d_1[0][0]
re_lu_1 (ReLU)	(None, 20, 64)	0	batch_normalization_1...
conv1d_2 (Conv1D)	(None, 20, 64)	12,352	re_lu_1[0][0]
batch_normalization_2 (BatchNormalization)	(None, 20, 64)	256	conv1d_2[0][0]
add (Add)	(None, 20, 64)	0	re_lu[0][0], batch_normalization_2...
re_lu_2 (ReLU)	(None, 20, 64)	0	add[0][0]
conv1d_3 (Conv1D)	(None, 20, 64)	12,352	re_lu_2[0][0]
batch_normalization_3 (BatchNormalization)	(None, 20, 64)	256	conv1d_3[0][0]
re_lu_3 (ReLU)	(None, 20, 64)	0	batch_normalization_3...
conv1d_4 (Conv1D)	(None, 20, 64)	12,352	re_lu_3[0][0]
batch_normalization_4 (BatchNormalization)	(None, 20, 64)	256	conv1d_4[0][0]
add_1 (Add)	(None, 20, 64)	0	re_lu_2[0][0], batch_normalization_4...
re_lu_4 (ReLU)	(None, 20, 64)	0	add_1[0][0]
global_average_pooling1d (GlobalAveragePooling1D)	(None, 64)	0	re_lu_4[0][0]
dense (Dense)	(None, 12)	780	global_average_poolin...

Training the model

```

import matplotlib.pyplot as plt
import seaborn as sns

sns.countplot(x=y_train)
plt.xticks(rotation=45)
plt.title("Class Distribution in Training Data")
plt.show()

```



```
import pandas as pd
from sklearn.utils import resample

X_train_df = pd.DataFrame(X_train, columns=[f"feature_{i}" for i in range(X_train.shape[1])])
y_train_df = pd.DataFrame(y_train, columns=["Workout"])

df = pd.concat([X_train_df, y_train_df], axis=1)

df_majority = df[df["Workout"] == 5]
df_minority = df[df["Workout"] != 5]

df_majority_downsampled = resample(df_majority,
                                   replace=False,
                                   n_samples=len(df_minority),
                                   random_state=42)

df_balanced = pd.concat([df_majority_downsampled, df_minority])

df_balanced = df_balanced.sample(frac=1, random_state=42)

X_train_balanced = df_balanced.drop(columns=["Workout"]).values
y_train_balanced = df_balanced["Workout"].values

import numpy as np
import pandas as pd
import tensorflow as tf
from sklearn.preprocessing import LabelEncoder

gpus = tf.config.list_physical_devices('GPU')
if gpus:
    try:
        for gpu in gpus:
            tf.config.experimental.set_memory_growth(gpu, True)
    except RuntimeError:
        pass

tf.keras.mixed_precision.set_global_policy('mixed_float16')

def load_and_preprocess_data(file_path):
    data = pd.read_csv(file_path)
    data.dropna(inplace=True)

    sensor_cols = ["A_x", "A_y", "A_z", "G_x", "G_y", "G_z", "C_1"]
    features = data[sensor_cols].values.astype(np.float32)

    le = LabelEncoder()
    labels = le.fit_transform(data["Workout"])

    return features, labels, len(le.classes_)

def build_fast_model(input_shape, num_classes):
    inputs = tf.keras.Input(shape=input_shape)
    x = tf.keras.layers.Conv1D(128, 3, padding='same')(inputs)
    x = tf.keras.layers.BatchNormalization()(x)
    x = tf.keras.layers.ReLU()(x)
```

```

x = tf.keras.layers.GlobalAveragePooling1D()(x)
outputs = tf.keras.layers.Dense(num_classes, activation='softmax', dtype='float32')(x)
return tf.keras.Model(inputs, outputs)

def train_pipeline():
    X, y, num_classes = load_and_preprocess_data("RecGym.csv")

    X = X.reshape(*X.shape, 1)

    train_ds = tf.data.Dataset.from_tensor_slices((X, y))
    train_ds = train_ds.shuffle(10000).batch(4096).prefetch(2)

    model = build_fast_model(X.shape[1:], num_classes)
    model.compile(
        optimizer=tf.keras.optimizers.AdamW(learning_rate=0.001, weight_decay=0.0001),
        loss='sparse_categorical_crossentropy',
        metrics=['accuracy']
    )

    history = model.fit(
        train_ds,
        epochs=200,
        verbose=1,
        callbacks=[
            tf.keras.callbacks.EarlyStopping(patience=3),
            tf.keras.callbacks.ReduceLROnPlateau(factor=0.5, patience=1)
        ]
    )

    return model

if __name__ == "__main__":
    model = train_pipeline()
    model.save("gym_activity_model.keras")

```

```

Epoch 1/200
1149/1149 ————— 261s 225ms/step - accuracy: 0.5328 - loss: 1.7045 - learning_rate: 0.0010
Epoch 2/200
/usr/local/lib/python3.11/dist-packages/keras/src/callbacks/early_stopping.py:153: UserWarning: Early stopping conditioned on metric `current` which has not been introduced yet: `current` is not found in the list of metrics: ['accuracy', 'loss']
  current = self.get_monitor_value(logs)
/usr/local/lib/python3.11/dist-packages/keras/src/callbacks/callback_list.py:145: UserWarning: Learning rate reduction is conditioned on metrics which have not been introduced yet: ['learning_rate']
  callback.on_epoch_end(epoch, logs)
1149/1149 ————— 256s 223ms/step - accuracy: 0.5664 - loss: 1.5336 - learning_rate: 0.0010
Epoch 3/200
1149/1149 ————— 254s 221ms/step - accuracy: 0.5686 - loss: 1.5130 - learning_rate: 0.0010
Epoch 4/200
1149/1149 ————— 257s 224ms/step - accuracy: 0.5700 - loss: 1.4977 - learning_rate: 0.0010
Epoch 5/200
1149/1149 ————— 257s 224ms/step - accuracy: 0.5714 - loss: 1.4851 - learning_rate: 0.0010
Epoch 6/200
1149/1149 ————— 265s 226ms/step - accuracy: 0.5730 - loss: 1.4723 - learning_rate: 0.0010
Epoch 7/200
1149/1149 ————— 263s 227ms/step - accuracy: 0.5740 - loss: 1.4630 - learning_rate: 0.0010
Epoch 8/200
1149/1149 ————— 264s 230ms/step - accuracy: 0.5752 - loss: 1.4547 - learning_rate: 0.0010
Epoch 9/200
1149/1149 ————— 260s 226ms/step - accuracy: 0.5760 - loss: 1.4477 - learning_rate: 0.0010
Epoch 10/200
1149/1149 ————— 261s 227ms/step - accuracy: 0.5770 - loss: 1.4416 - learning_rate: 0.0010
Epoch 11/200
1149/1149 ————— 260s 226ms/step - accuracy: 0.5777 - loss: 1.4356 - learning_rate: 0.0010
Epoch 12/200
1149/1149 ————— 261s 225ms/step - accuracy: 0.5784 - loss: 1.4306 - learning_rate: 0.0010
Epoch 13/200
1149/1149 ————— 266s 229ms/step - accuracy: 0.5794 - loss: 1.4245 - learning_rate: 0.0010
Epoch 14/200
1149/1149 ————— 319s 227ms/step - accuracy: 0.5799 - loss: 1.4197 - learning_rate: 0.0010
Epoch 15/200
1149/1149 ————— 261s 227ms/step - accuracy: 0.5807 - loss: 1.4152 - learning_rate: 0.0010
Epoch 16/200
1149/1149 ————— 263s 229ms/step - accuracy: 0.5811 - loss: 1.4121 - learning_rate: 0.0010
Epoch 17/200
1149/1149 ————— 316s 224ms/step - accuracy: 0.5820 - loss: 1.4085 - learning_rate: 0.0010
Epoch 18/200
1149/1149 ————— 274s 238ms/step - accuracy: 0.5823 - loss: 1.4058 - learning_rate: 0.0010
Epoch 19/200
1149/1149 ————— 268s 233ms/step - accuracy: 0.5830 - loss: 1.4012 - learning_rate: 0.0010
Epoch 20/200
1149/1149 ————— 320s 231ms/step - accuracy: 0.5835 - loss: 1.3977 - learning_rate: 0.0010
Epoch 21/200
1149/1149 ————— 323s 232ms/step - accuracy: 0.5841 - loss: 1.3946 - learning_rate: 0.0010
Epoch 22/200
1149/1149 ————— 318s 229ms/step - accuracy: 0.5841 - loss: 1.3924 - learning_rate: 0.0010
Epoch 23/200
1149/1149 ————— 325s 231ms/step - accuracy: 0.5851 - loss: 1.3898 - learning_rate: 0.0010

```

Epoch 24/200

1149/1149 ————— 320s 230ms/step - accuracy: 0.5852 - loss: 1.3874 - learning_rate: 0.0010

Epoch 25/200

1149/1149 ————— 263s 229ms/step - accuracy: 0.5853 - loss: 1.3854 - learning_rate: 0.0010

Epoch 26/200

1149/1149 ————— 265s 230ms/step - accuracy: 0.5858 - loss: 1.3838 - learning_rate: 0.0010